

Improve fairness of AI systems

Learn about AI fairness from our guides and use cases. Assess and mitigate fairness issues using our Python toolkit. Join our community and contribute metrics, algorithms, and other resources.

Get Started




Fairlearn is an open-source, community-driven project to help data scientists improve fairness of AI systems.

FAIRNESS IS SOCIOTECHNICAL

Fairness of AI systems is about more than simply running lines of code. In each use case, both societal and technical aspects shape who might be harmed by AI systems and how. There are many complex sources of unfairness and a variety of societal and technical processes for mitigation, not just the mitigation algorithms in our library.

Throughout this website, you can find resources on how to think about fairness as sociotechnical, and how to use Fairlearn's metrics and algorithms while considering the AI system's broader societal context.



USE CASE | CREDIT-CARD LOANS

Assessment and mitigation of fairness issues in credit-card default models



When making a decision to approve or decline a loan, financial services organizations use a variety of models, including a model that predicts the applicant's probability of default. These predictions are sometimes used to automatically reject or accept an application, directly impacting both the applicant and the organization.

In this scenario, fairness-related harms may arise when the model makes more mistakes for some groups of applicants compared to others. We use Fairlearn to assess how different groups, defined in terms of their sex, are affected and how the observed disparities may be mitigated.

Jupyter Notebook



Get Started with Fairlearn


To get started, install the Fairlearn package. But the process does not end there! See our user guide and other resources to understand what fairness means for your use case.

If you run into any issues, reach out on [Gitter](#).

Fairlearn can be installed with pip from PyPI as follows:

```
pip install fairlearn
```

Quickstart




Resources




User Guide

Learn more about fairness in AI, fairness metrics, and mitigation algorithms.




API Documentation

Library reference with examples.



Contributor Guide

Help us with case studies, documentation, or code. There are many ways to contribute, regardless of your background or expertise.







Quickstart

Cursus in hac habitasse platea dictumst. Blandit turpis cursus in hac habitasse platea dictumst quisque.

Join the Community

Fairlearn community consists of open source contributors, data science practitioners, and responsible AI enthusiasts across many disciplines and locations. Join the latest conversation, ask usage questions, learn about the most recent updates, and find out how to get involved.

 Gitter Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore reprehenderit.	 Stack Overflow Consequat aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu voluptate.	 GitHub Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore reprehenderit.	 Twitter Consequat aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu voluptate.
--	--	--	---



Contribute to Fairlearn

Fairlearn is built and maintained by open source contributors with a variety of backgrounds and expertise. Join the effort and contribute feedback, metrics, algorithms, visualizations, ideas and more, so we can evolve the toolkit together!

Contributor Guide

Resources

- Quickstart
- User Guide
- API Documentation
- Examples
- Contributor Guide
- Discussions on GitHub
- Conversations on Gitter

Fairlearn

- About Fairlearn
- Code of Conduct
- Governance
- Project Roadmap
- FAQ
- Feedback
- Contact Us

Stay Connected

Be the first to hear about the latest Fairlearn announcements, new publications, and events.

email@gmail.com

