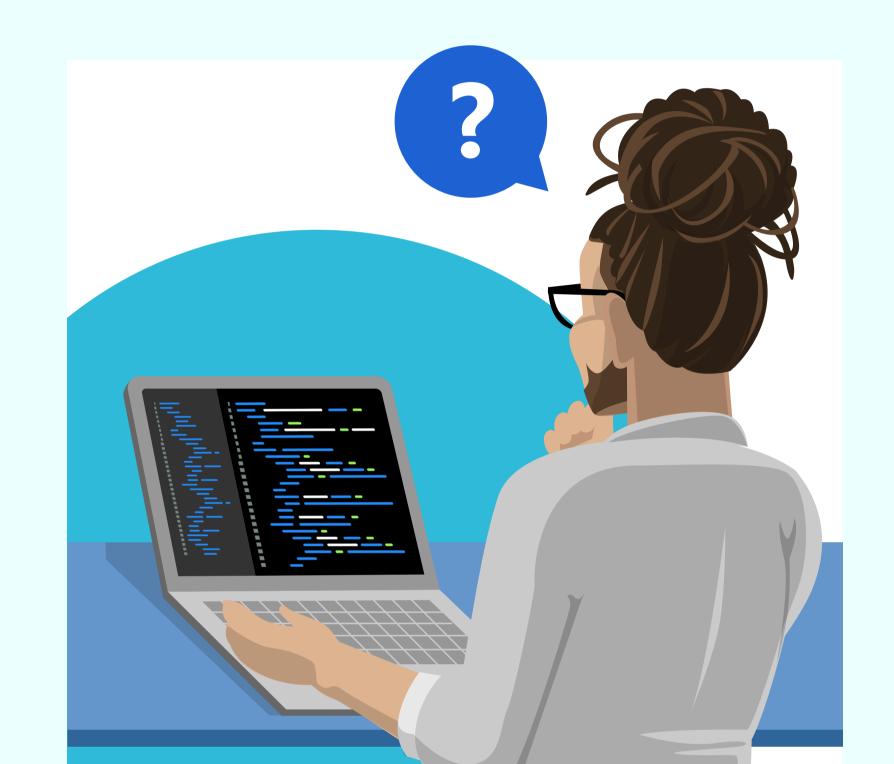


Fairness of AI systems is about more than simply running lines of code. In each use case, both societal and technical aspects shape who might be harmed by Al systems and how. There are many complex sources of unfairness and a variety of societal and technical processes for mitigation, not just the mitigation algorithms in our library. Throughout this website, you can find resources on how to think about fairness as sociotechnical, and how to use Fairlearn's metrics and algorithms while considering the AI system's broader societal context.



USE CASE | CREDIT-CARD LOANS Assessment and mitigation of fairness issues in credit-card default models



When making a decision to approve or decline a loan, financial services organizations use a variety of models, including a model that predicts the applicant's probability of default. These predictions are sometimes used to automatically reject or accept an application, directly impacting both the applicant and the organization.

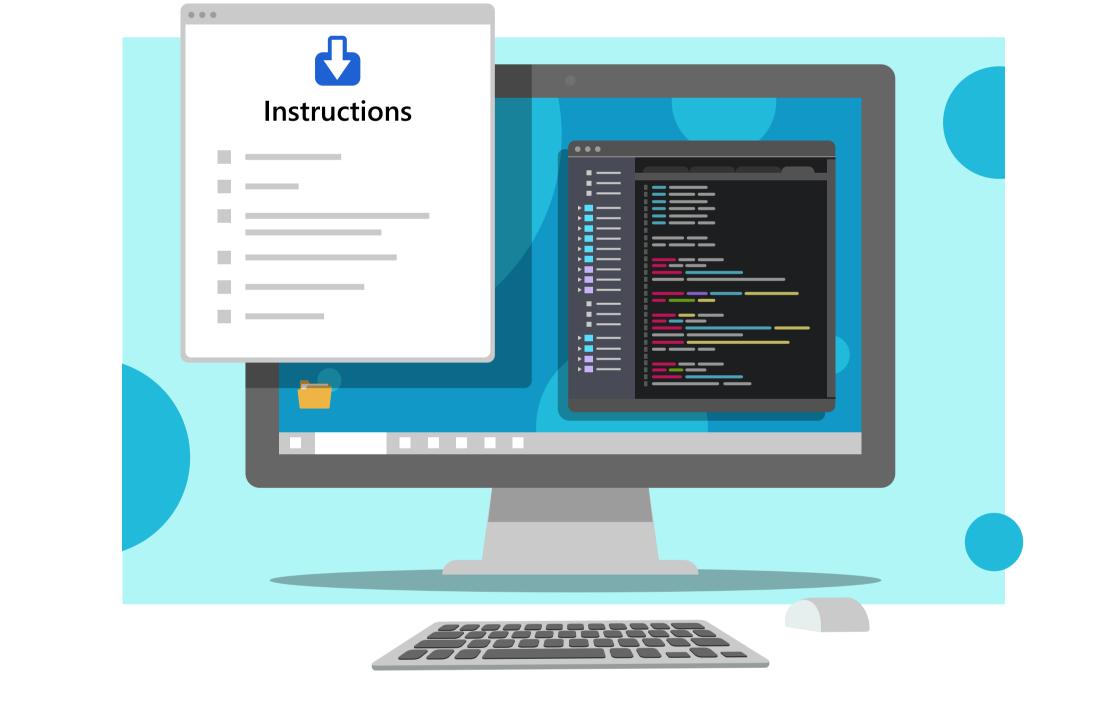
In this scenario, fairness-related harms may arise when the model makes more mistakes for some groups of applicants compared to others. We use Fairlearn to assess how different groups, defined in terms of their sex, are affected and how the observed disparities may be mitigated.

Get Started with Fairlearn

To get started, install the Fairlearn package. But the process does not end there! See our user guide and other resources to understand what fairness means for your use case.

If you run into any issues, reach out on Gitter.

Fairlearn can be installed with pip from PyPI as follows: pip install fairlearn



Resources

User Guide Learn more about fairness in Al, fairness metrics, and mitigation **API** Documentation

Library reference with examples.

Contributor Guide Help us with case studies, many ways to contribute, regardless of your background or expertise.

Quickstart Cursus in hac habitasse platea dictumst. Blandit turpis cursus in hac habitasse platea dictumst quisque.

Join the Community

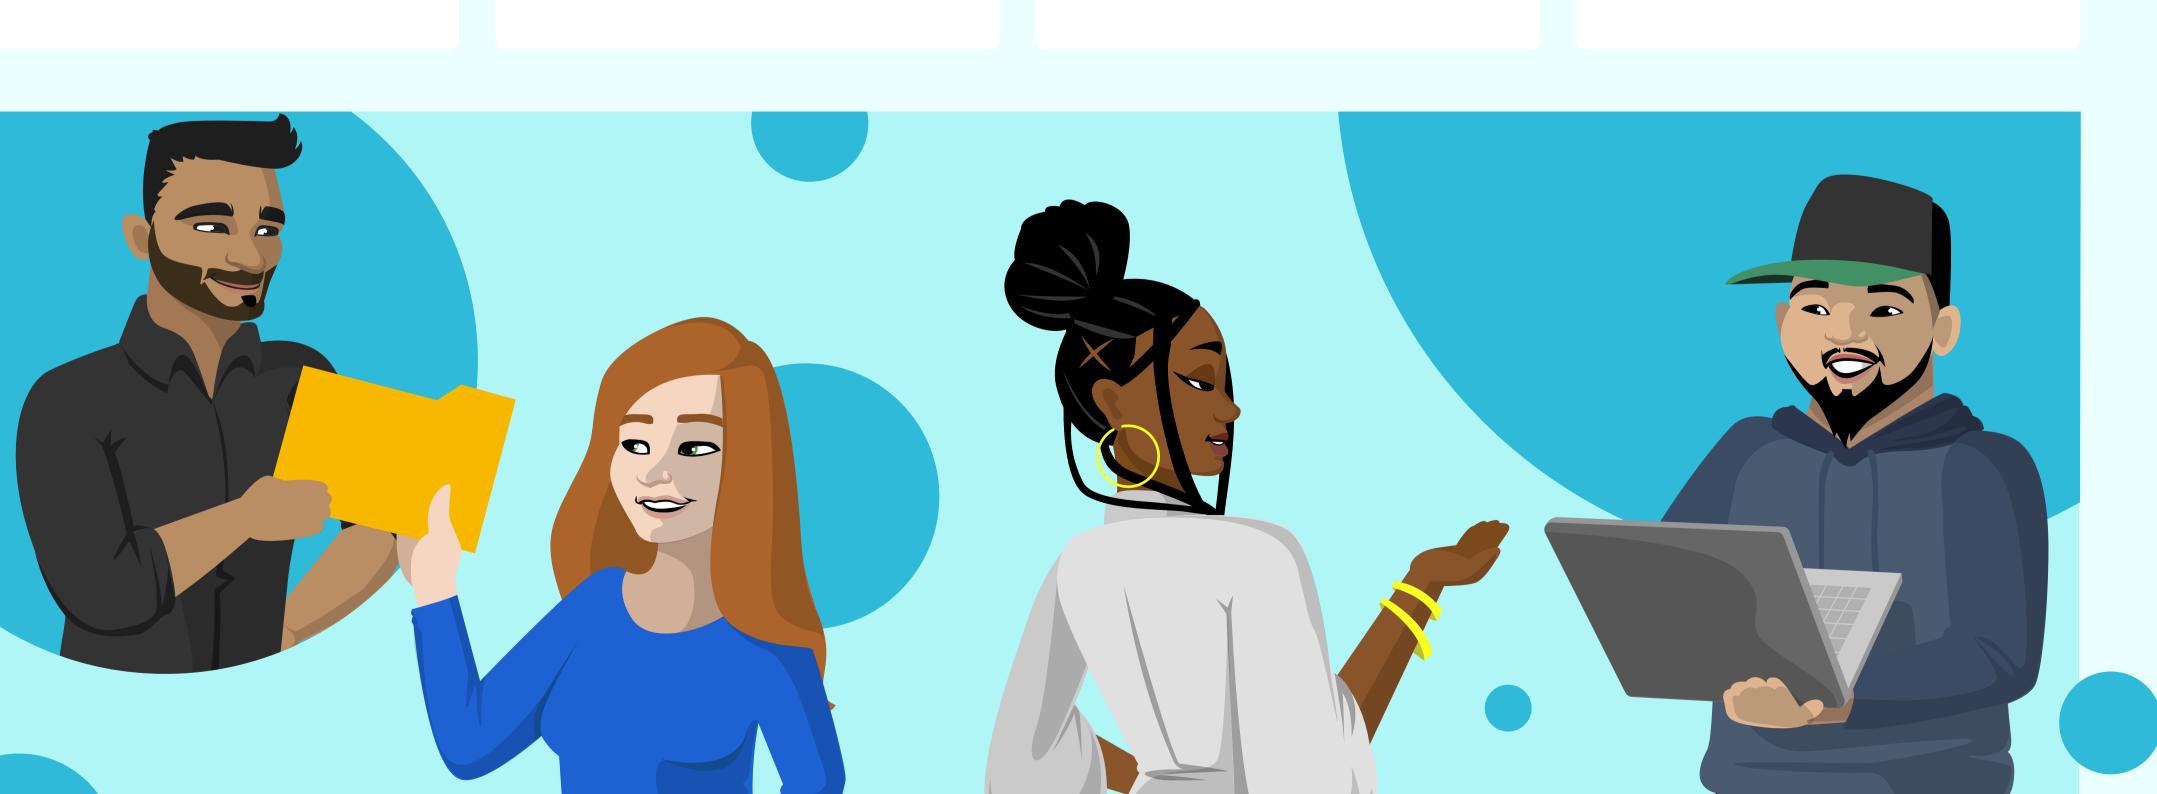
Fairlearn community consists of open source contributors, data science practitioners, and responsible AI enthusiasts across many disciplines and locations. Join the latest conversation, ask usage questions, learn about the most recent updates, and find out how to get involved.

□ Gitter Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore reprehenderit.

Stack Overflow Consequat aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu voluptate.

GitHub Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore reprehenderit

Twitter Consequat aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu voluptate.



Contribute to Fairlearn

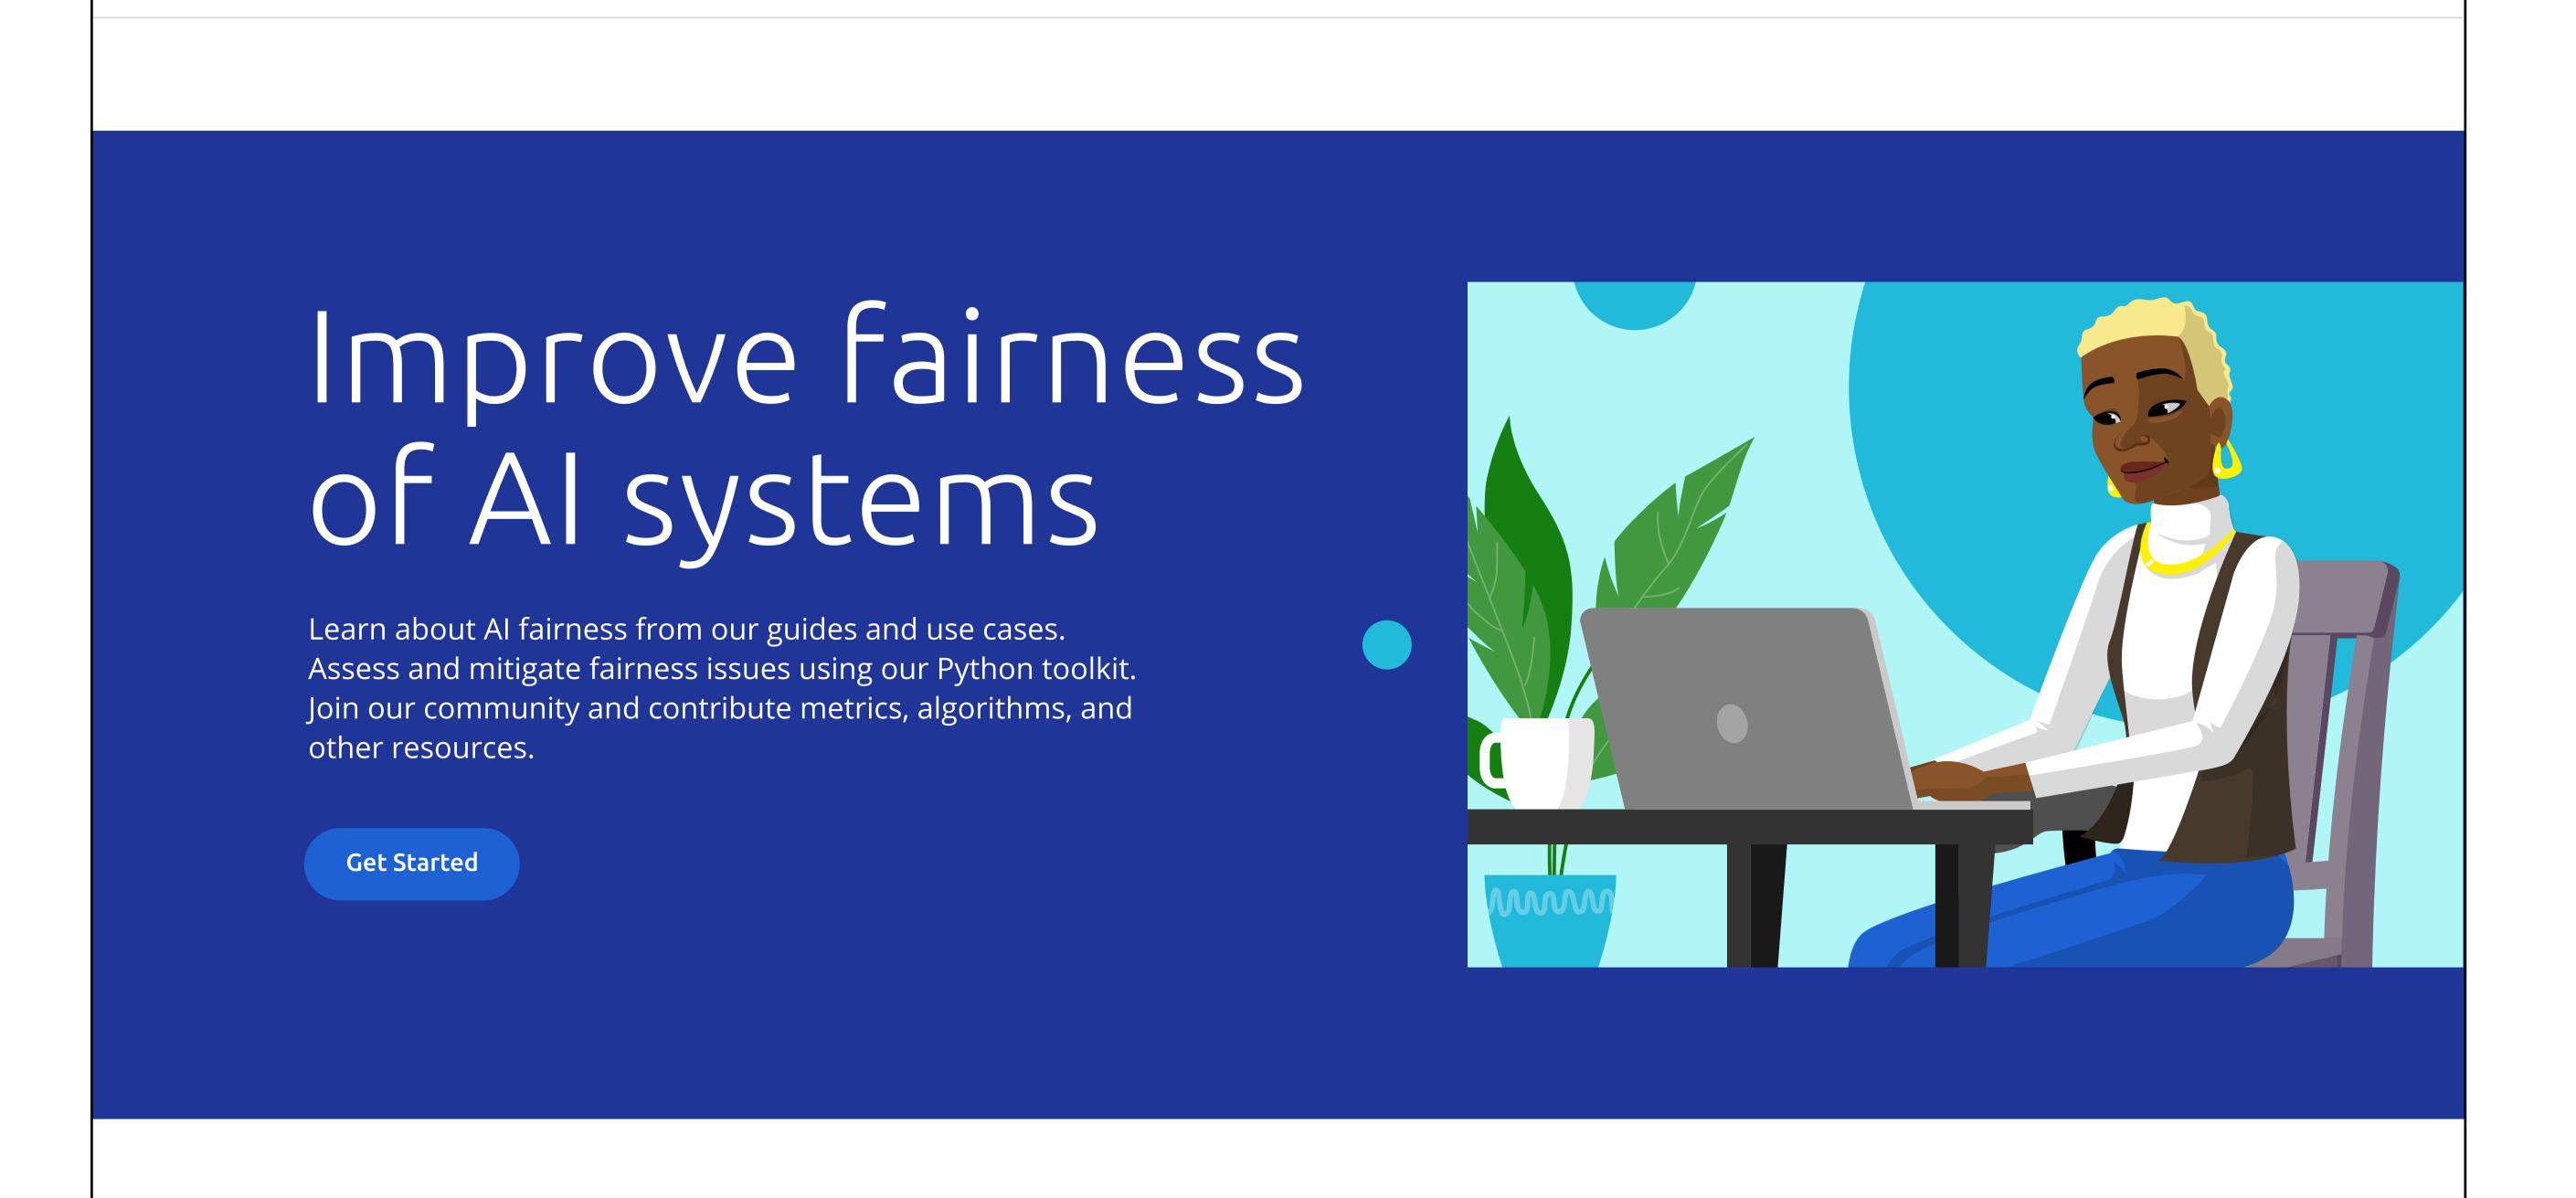
Fairlearn is built and maintained by open source contributors with a variety of backgrounds and expertise. Join the effort and contribute feedback, metrics, algorithms, visualizations, ideas and more, so we can evolve the toolkit together!

Contributor Guide

#fairlearn Fairlearn Resources Stay Connected About Fairlearn User Guide Code of Conduct **API** Documentation Governance Project Roadmap Contact Us Conversations on Gitter **Fairlearn** Terms of Use Privacy © 2021

#1F3598

Darkest Fairlearn blue



Fairlearn is an open-source, community-driven project to help data scientists improve fairness of AI systems.

FAIRNESS IS SOCIOTECHNICAL Fairness of AI systems is about more than simply running lines of code. In each use case, both societal and technical aspects shape who might be harmed by Al systems and how. There are many complex sources of unfairness and a variety of societal and technical processes for mitigation, not just the mitigation algorithms in our library. Throughout this website, you can find resources on how to think about fairness as sociotechnical, and

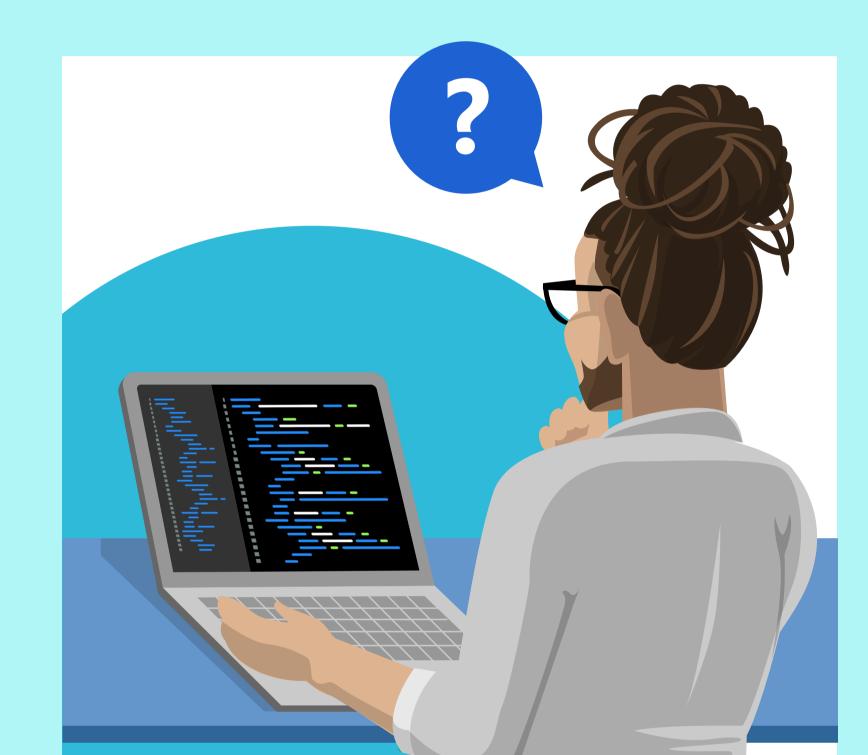
how to use Fairlearn's metrics and algorithms while

considering the AI system's broader societal context.



GitHub Q

USE CASE | CREDIT-CARD LOANS Assessment and mitigation of fairness issues in credit-card default models



When making a decision to approve or decline a loan, financial services organizations use a variety of models, including a model that predicts the applicant's probability of default. These predictions are sometimes used to automatically reject or accept an application, directly impacting both the applicant and the organization. In this scenario, fairness-related harms may arise when the model makes more mistakes for some groups of applicants compared to others. We use Fairlearn to assess how different groups, defined in terms of their sex, are

affected and how the observed disparities may be mitigated.

Jupyter Notebook

Get Started with Fairlearn

To get started, install the Fairlearn package. But the process does not end there! See our user guide and other resources to understand what fairness means for your use case. If you run into any issues, reach out on Gitter.

Fairlearn can be installed with pip from PyPI as follows: pip install fairlearn

Resources

User Guide

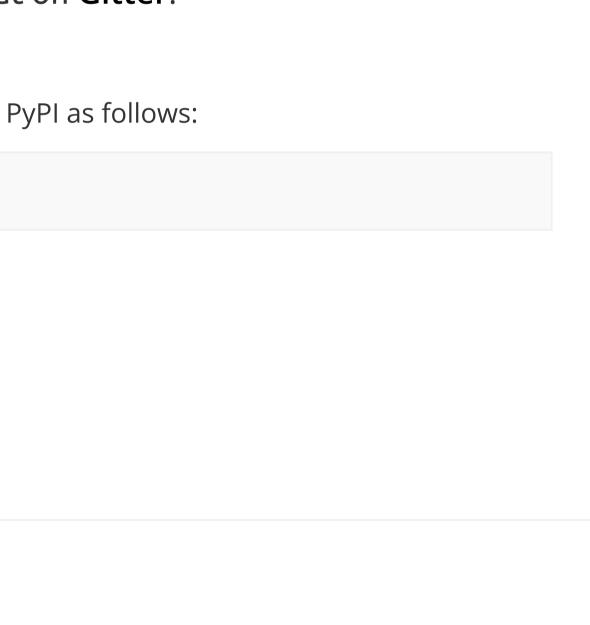
algorithms.

□ Gitter

Duis aute irure dolor in

Learn more about fairness in Al

fairness metrics, and mitigation



API Documentation

Library reference with examples.

Stack Overflow

Contributor Guide Help us with case studies, documentation, or code. There are

Quickstart Cursus in hac habitasse platea dictumst. Blandit turpis cursus in hac habitasse platea dictumst quisque. of your background or expertise.

Join the Community Fairlearn community consists of open source contributors, data science practitioners,

and responsible AI enthusiasts across many disciplines and locations. Join the latest conversation, ask usage questions, learn about the most recent updates, and find out how to get involved.

reprehenderit in voluptate velit reprehenderit in voluptate velit esse cillum dolore reprehenderit esse cillum dolore eu voluptate.

GitHub Consequat aute irure dolor in Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore reprehenderit.

Twitter Consequat aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu voluptate.

Contribute to Fairlearn

Fairlearn is built and maintained by open source contributors with a variety of backgrounds and expertise. Join the effort and contribute feedback, metrics, algorithms, visualizations, ideas and more, so we can evolve the toolkit together! Contributor Guide

Fairlearn Resources Quickstart About Fairlearn User Guide Code of Conduct **API** Documentation Governance Examples Project Roadmap Contributor Guide Contact Us Conversations on Gitter

Fairlearn Terms of Use Privacy

Stay Connected Be the first to hear about the latest Fairlearn announcements, new publications, and email@gmail.com

#fairlearn

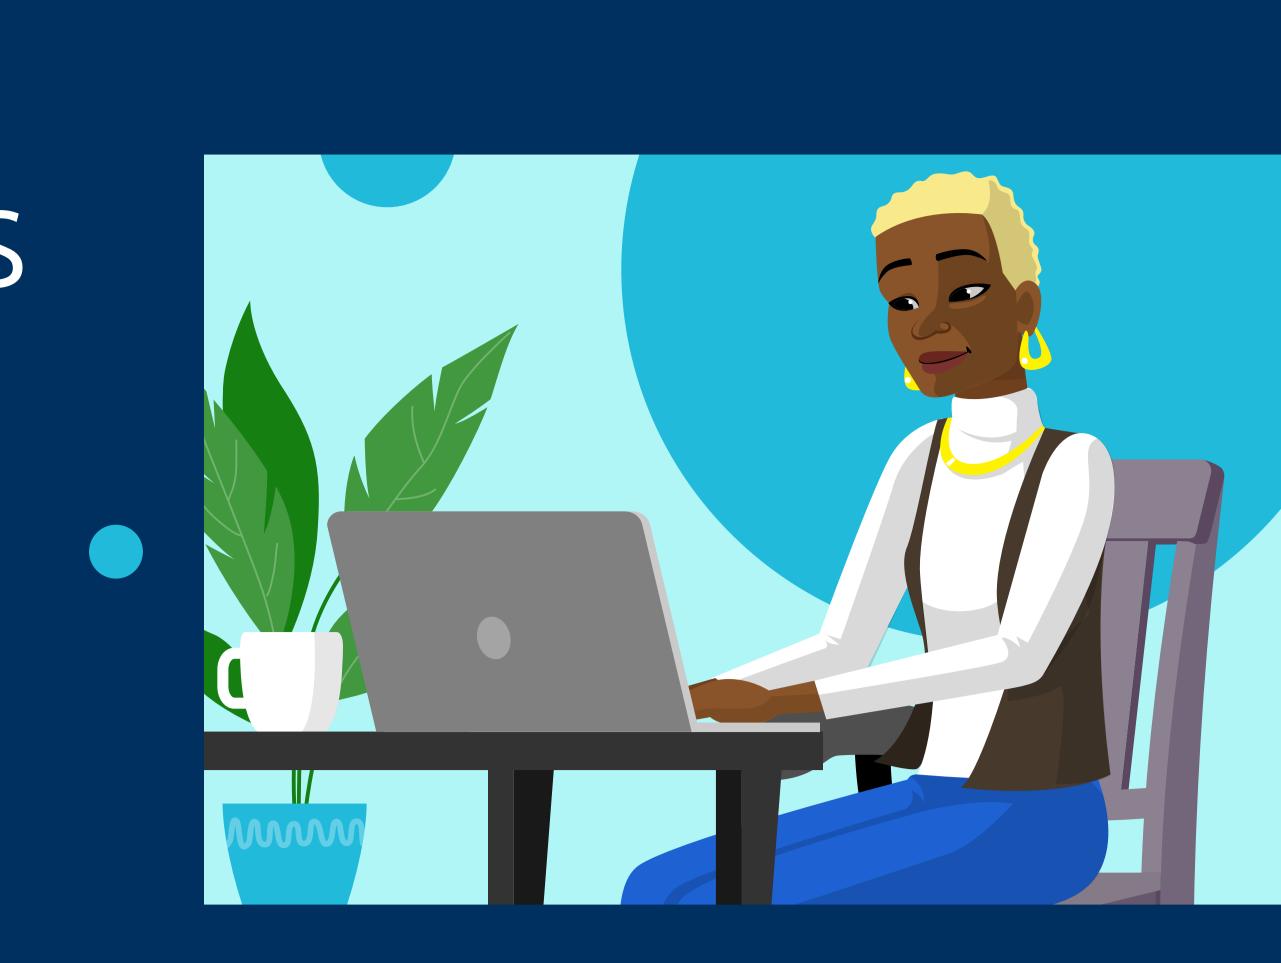
© 2021

#003060

Navy blue



Learn about Al fairness from our guides and use cases. Assess and mitigate fairness issues using our Python toolkit. Join our community and contribute metrics, algorithms, and other resources.



GitHub Q

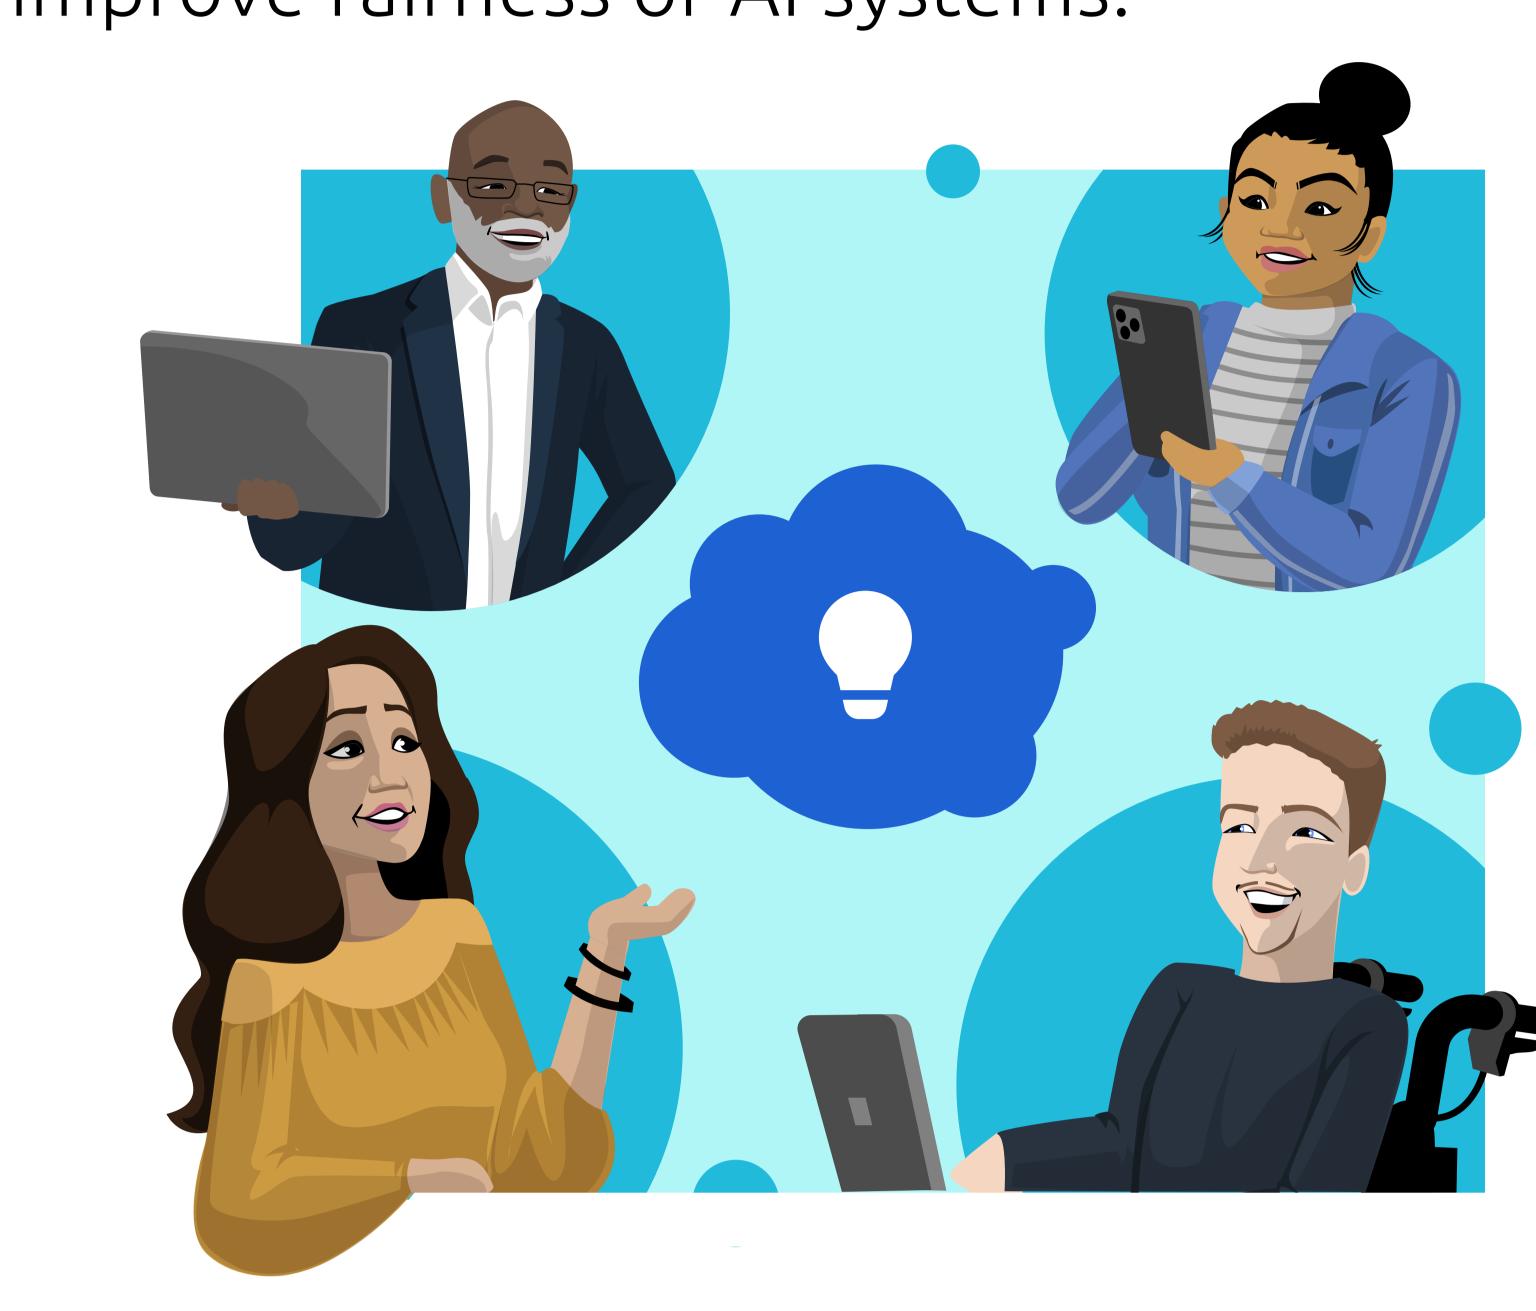
Fairlearn is an open-source, community-driven project to help data scientists improve fairness of AI systems.

FAIRNESS IS SOCIOTECHNICAL Fairness of AI systems is about more than simply running lines of code. In each use case, both societal and technical aspects shape who might be harmed by Al systems and how. There are many complex sources of unfairness and a variety of societal and technical

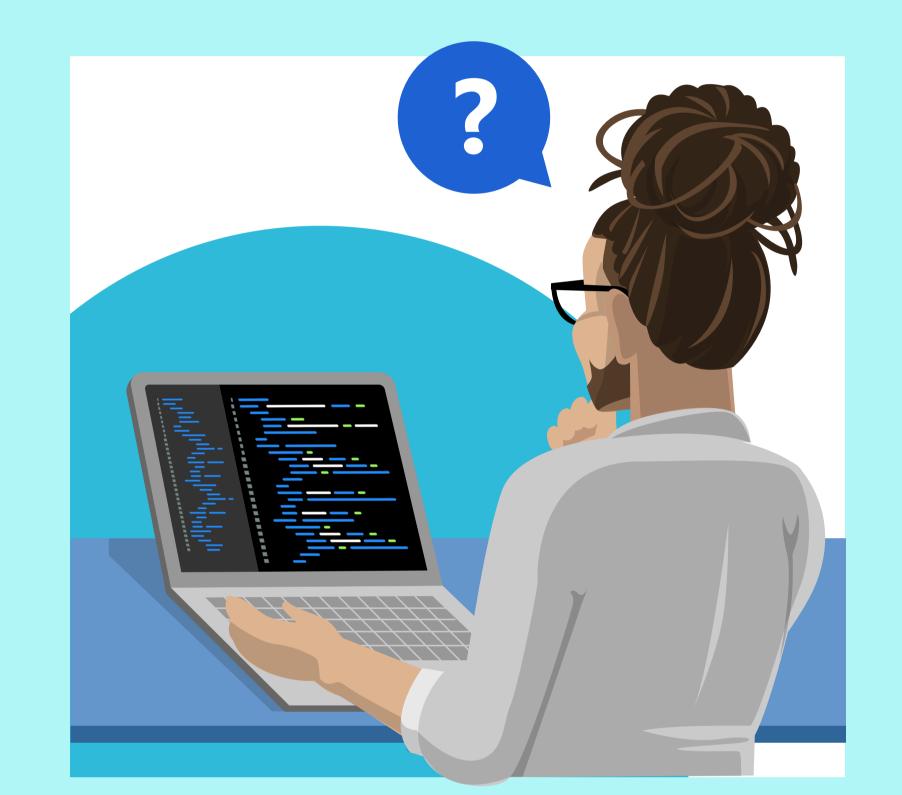
algorithms in our library.

Throughout this website, you can find resources on how to think about fairness as sociotechnical, and how to use Fairlearn's metrics and algorithms while considering the AI system's broader societal context.

processes for mitigation, not just the mitigation



USE CASE | CREDIT-CARD LOANS Assessment and mitigation of fairness issues in credit-card default models



When making a decision to approve or decline a loan, financial services organizations use a variety of models, including a model that predicts the applicant's probability of default. These predictions are sometimes used to automatically reject or accept an application, directly impacting both the applicant and the organization.

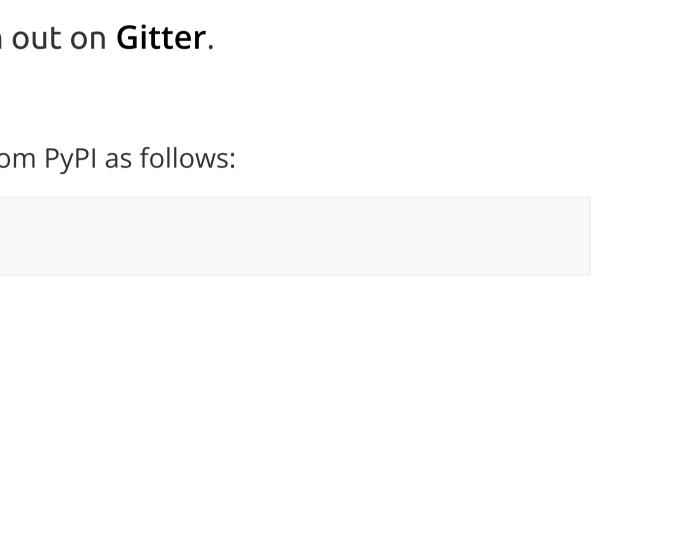
In this scenario, fairness-related harms may arise when the model makes more mistakes for some groups of applicants compared to others. We use Fairlearn to assess how different groups, defined in terms of their sex, are affected and how the observed disparities may be mitigated.

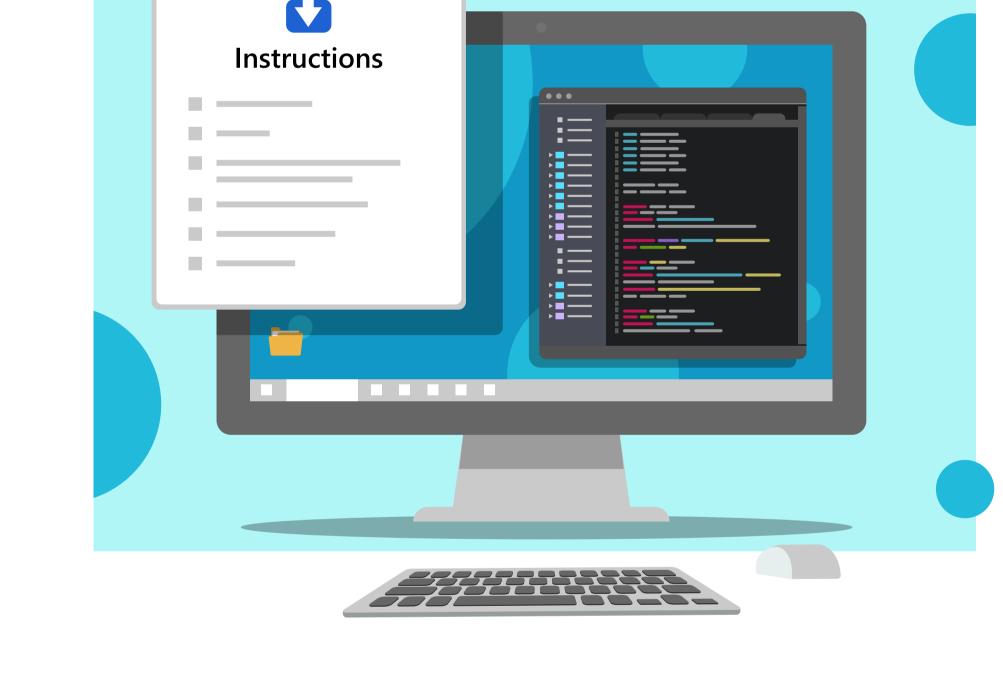
Jupyter Notebook

Get Started with Fairlearn

To get started, install the Fairlearn package. But the process does not end there! See our user guide and other resources to understand what fairness means for your use case. If you run into any issues, reach out on Gitter.

Fairlearn can be installed with pip from PyPI as follows: pip install fairlearn





Resources

User Guide Learn more about fairness in A fairness metrics, and mitigation

API Documentation Library reference with examples.

Contributor Guide Help us with case studies, documentation, or code. There are many ways to contribute, regardless of your background or expertise.

Quickstart Cursus in hac habitasse platea dictumst. Blandit turpis cursus in hac habitasse platea dictumst quisque.

Join the Community Fairlearn community consists of open source contributors, data science practitioners,

and responsible Al enthusiasts across many disciplines and locations. Join the latest conversation, ask usage questions, learn about the most recent updates, and find out how to get involved.

Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore reprehenderit.

Fairlearn Terms of Use Privacy

□□ Gitter

Consequat aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu voluptate.

GitHub Stack Overflow

Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore reprehenderit

Consequat aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu voluptate.

Twitter



Contribute to Fairlearn Fairlearn is built and maintained by open source contributors with a variety of backgrounds and expertise. Join the effort and contribute feedback, metrics, algorithms, visualizations, ideas and more, so we can evolve the toolkit together! Contributor Guide

#fairlearn Stay Connected Fairlearn Resources Quickstart About Fairlearn User Guide Code of Conduct API Documentation Governance Project Roadmap Examples Contact Us Conversations on Gitter